

Malhar Sham Ghogare

ghogare.m@northeastern.edu

Shakespeare QA Bot Using Hugging Face Transformers

1. Data Processing

Code:

```
# Step 1: Extract text from the PDF
def extract_text_from_pdf(pdf_path):
    reader = PdfReader(pdf_path)
    text = ""
    for page in reader.pages:
        text += page.extract_text()
    return text
```

Purpose:

- This function extracts raw text from a PDF file.
- It reads each page of the PDF using the PdfReader from the PyPDF2 library and concatenates the text into a single string.

Usage:

- Provide the path to the PDF file as an argument to this function.

1.2. Cleaning Extracted Text

Code:

```
def clean_text(raw_text):
    # Basic cleaning: remove headers, footers, and excessive whitespace
    text = re.sub(r"\n{2,}", "\n", raw_text) # Replace multiple newlines with a single
    text = re.sub(r"\s{2,}", " ", text) # Replace multiple spaces with a single space
    text = re.sub(r"\b\d+\b", "", text) # Remove isolated numbers (like page numbers)
    # Retain only paragraphs with valid content
    paragraphs = [
        para.strip() for para in text.split("\n")
        if len(para.strip()) > 20 and re.search(r"^[A-Za-z]", para)
    ]
    return " ".join(paragraphs)

    # Step 3: Remove irrelevant sections using stricter filtering
    paragraphs = [
        para for para in paragraphs
        if not re.search(r"(COPYRIGHT|PROJECT GUTENBERG|ELECTRONIC VERSION|SERVICE|DOWNLOADED BY)", para)
        and not re.search(r"^[s]*(DISTRIBUTED|PERSONAL USE ONLY|COMMERCIALY)", para, re.IGNORECASE)
        and len(para.strip()) > 20 # Exclude short lines likely to be non-content
    ]
    return " ".join(paragraphs)

    # Step 4: Retain only paragraphs with recognizable Shakespearean content
    paragraphs = [
        para for para in paragraphs
        if re.search(r"^[A-Za-z]", para) # Starts with letters
    ]

    return " ".join(paragraphs)
```

Output:

- Returns cleaned text as a single string, ready for further processing.

Usage:

- Pass the `raw_text` (output from `extract_text_from_pdf`) to this function.

1.3. Segmenting the Text

Code:

```
def segment_text(text, segment_size=500):  
    return [text[i:i + segment_size] for i in range(0, len(text), segment_size)]
```

Purpose:

- Divides the cleaned text into smaller, fixed-sized chunks for efficient processing.
- The default segment size is 500 characters.

Output:

- A list of text segments.

Usage:

- Pass the `cleaned_text` to this function.

1.4. Extracting Named Entities

Code:

```
def extract_entities(text, nlp):  
    doc = nlp(text)  
    return {ent.text.lower() for ent in doc.ents}
```

Purpose:

- Uses spaCy's Named Entity Recognition (NER) to identify and extract named entities (e.g., people, places, dates) from a text segment.

Output:

- A set of named entities extracted from the text.

Usage:

- Load spaCy's English model and pass it along with a text segment to this function.

1.5. Preprocessing Segments with Named Entities

Code:

```
def preprocess_segments_with_entities(segments, nlp):  
    # Cache entities for all segments  
    segment_entities = []  
    for segment in segments:  
        entities = extract_entities(segment, nlp)  
        segment_entities.append((segment, entities))  
    return segment_entities
```

Purpose:

- Extracts and caches named entities for each text segment.
- Stores the segment and its corresponding entities as a tuple for later use.

Output:

- A list of tuples, where each tuple contains:
 1. A text segment.
 2. A set of named entities extracted from the segment.

Usage:

- Pass the segmented text and spaCy model to this function.

1.6. Finding Relevant Segments

Code:

```
# Step 5: Find Relevant Segment Using Preprocessed Entities
def find_relevant_segment(question, segment_entities, nlp):
    question_doc = nlp(question)
    question_entities = {ent.text.lower() for ent in question_doc.ents}

    best_segment = ""
    max_matches = 0
    for segment, entities in segment_entities:
        matches = len(question_entities.intersection(entities))
        if matches > max_matches:
            max_matches = matches
            best_segment = segment
    return best_segment
```

Purpose:

- Matches named entities in the user's question with those in the preprocessed segments.
- Selects the segment with the highest number of matching entities as the most relevant context.

Output:

- The most relevant text segment.

Usage:

- Pass the question, preprocessed segments (segment_entities), and spaCy model to this function.

2. Model Selection and Setup

Code:

```
# Step 3: Set up the Hugging Face Q&A pipeline
def setup_pipeline():
    return pipeline("question-answering", model="distilbert-base-uncased", tokenizer="distilbert-base-uncased")

# Step 4: Ask questions and get answers
def ask_question(qa_pipeline, context, question):
    result = qa_pipeline(question=question, context=context)
    return result[0]["answer"]
```

Purpose:

- Initializes the Hugging Face Question-Answering pipeline using the pre-trained distilbert-base-uncased model.

Model Description:

This Question and Answer bot leverages **DistilBERT**, a smaller, faster, and lighter variant of the BERT (Bidirectional Encoder Representations from Transformers) model, fine-tuned for question-answering tasks.

Key Features of the Model:

- **Transformer Architecture:**
 - DistilBERT is based on the Transformer architecture, which excels in understanding contextual relationships between words and phrases in text.
- **Pretrained on Large Text Datasets:**
 - The model has been pre-trained on extensive corpora (e.g., Wikipedia and BookCorpus), making it capable of understanding a wide variety of natural language queries.
- **Fine-Tuned for Question Answering:**
 - Fine-tuning on datasets like SQuAD (Stanford Question Answering Dataset) enables the model to extract precise answers from a given context.

Why Use DistilBERT?

- **Efficiency:** DistilBERT is 60% smaller than BERT while retaining 97% of its language understanding capabilities. This makes it ideal for real-time applications like this bot.
- **Speed:** It runs faster than standard BERT models, making the QCA experience more seamless for users.
- **Accuracy:** Its fine-tuning on question-answering tasks ensures high-quality answers when provided with a relevant context.

How It Works in This Bot:

1. **Input Question:**
 - The user enters a question about Shakespeare's works.

2. Context Extraction:

- The bot first identifies the most relevant text segment from Shakespeare's corpus using Named Entity Recognition (NER) and entity matching techniques.

3. Answer Generation:

- The question and extracted context are passed to the DistilBERT model, which generates a concise and accurate answer.

Output:

- A pipeline object for question-answering tasks.

Usage:

- Call this function to set up the pipeline.
- Example:

3. Design Q A system

Code:

```
# Main Q&A Function
def shakespeare_qa(question):

    try:
        answer = qa_pipeline(question=question, context=relevant_segment)
        return answer['answer']
    except Exception as e:
        return f"Sorry, I couldn't find an answer. Error: {str(e)}"
```

Purpose:

- Passes the user's question and the relevant context to the Hugging Face pipeline.
- Retrieves and returns the answer from the pipeline's output.

Usage:

- Provide the pipeline object, a text segment (context), and a question to this function.

4. Implementing User Interface

Code:

```
import gradio as gr

# Create Gradio Interface
iface = gr.Interface(
    fn=shakespeare_qa,
    inputs=gr.Textbox(label="Ask a question about Shakespeare's works"),
    outputs=gr.Textbox(label="Answer"),
    title="Shakespeare Works Q&A",
    description="Ask questions about characters, plots, and themes in Shakespeare's works.",
    examples=[
        "Who is Bertram?",
        "What happens in Romeo and Juliet?",
        "Describe Macbeth's character",
        "Who wrote these plays?"
    ]
)

# Launch the interface
iface.launch(share=True)
```

Purpose:

- This code sets up an interactive web interface for the QCA system using **Gradio**.
- Users can input questions related to Shakespeare's works and receive answers in real time.

Usage:

- **gr.Interface:**
 - **fn:** Specifies the function to handle user queries. In this case, `shakespeare_qa` is the function that processes questions and retrieves answers.
 - **inputs:** Defines the input field for the interface. Here, it's a text box labeled *"Ask a question about Shakespeare's works"*.
 - **outputs:** Defines the output field for the interface. Here, it's another text box labeled *"Answer"*.

- **title:** Sets the title of the interface, displayed at the top of the page.
- **description:** Provides additional details about the interface's functionality.
- **examples:** Displays pre-defined example questions for users to try.
- **Launching the Interface:**
 - **iface.launch(share=True):**
 - Starts the Gradio web interface.
 - The share=True option generates a public link, allowing anyone to access the interface via a web browser.

How to Use the Web Interface:

1. Launching the Interface:

- Run the code block in your Python environment.
- Gradio will generate a local URL (e.g., `http://127.0.0.1:7860`) and a public URL if `share=True` is used (e.g., `https://1234abcd.gradio.app`).

2. Accessing the Interface:

- Open the public URL in a web browser.
- The interface will display the following components:
 - **Title:** *"Shakespeare Works Q&A"*
 - **Description:** *"Ask questions about characters, plots, and themes in Shakespeare's works."*
 - **Input Textbox:** A field to type your question.
 - **Output Textbox:** A field to display the answer.
 - **Examples:** Predefined example questions to help users get started.

3. Asking Questions:

- Type a question into the input text box (e.g., *"Who is Bertram?"*).
- Click the *"Submit"* button (or press Enter).

- The interface will process the question using the QCA pipeline and display the answer in the output text box.

4. Using Predefined Examples:

- Click on any example question (e.g., "*What happens in Romeo and Juliet?*").
- The interface will automatically populate the input box with the selected question.

Instructions to Set Up and Execute the Code:

Prerequisites

1. Install the required libraries:

```
import re
from transformers import pipeline
from PyPDF2 import PdfReader
import spacy
import gradio as gr
```

Steps to Run the Code

1. Extract and Clean Text:

- Use `extract_text_from_pdf` to read the PDF.
- Clean the text using `clean_text`.

2. Preprocess Segments:

- Segment the text using `segment_text`.
- Extract and cache entities with `preprocess_segments_with_entities`.

3. Set Up the Pipeline:

- Initialize the QCA pipeline with `setup_pipeline`.

4. Ask Questions:

- Find the relevant segment using `find_relevant_segment`.
- Retrieve the answer using `ask_question`.

5. Retrieve Answer:

- Use the ask_question function to pass the question and the relevant segment to the QCA pipeline. Clean the text using clean_text.
- The pipeline will extract and return the answer.

6. Activate User Interface:

Set Up Gradio Interface:

- The Gradio code (as shown in the image) sets up a user-friendly web interface.
- **Key Features:**
 - **fn=shakespeare_qa:** Connects the question-answering function (shakespeare_qa) to the interface.
 - **Input:** A text box labeled "Ask a question about Shakespeare's works".
 - **Output:** A text box labeled "Answer" for displaying the generated answers.
 - **Examples:** Predefined sample questions like:
 - "Who is Bertram?"
 - "What happens in Romeo and Juliet?"
 - "Describe Macbeth's character"
 - "Who wrote these plays?"

Launch the Interface:

- Use the iface.launch(share=True) function to start the Gradio interface.
- This generates a public URL that can be shared for accessing the QCA bot

The screenshot shows a web interface titled "Shakespeare Works Q&A". Below the title is a subtitle: "Ask questions about characters, plots, and themes in Shakespeare's works." The interface is divided into two main sections. On the left, there is a text input box with the placeholder "Ask a question about Shakespeare's works" and a sample question "Who is Bertram?". Below this input box are three buttons: "Clear", "Submit" (highlighted in orange), and "Flag". On the right, there is a text output box with the placeholder "Answer" and a sample answer "and Servant to the Countess of Rousillon". At the bottom of the interface, there is a section labeled "Examples" with four buttons: "Who is Bertram?", "What happens in Romeo and Juliet?", "Describe Macbeth's character", and "Who wrote these plays?".